

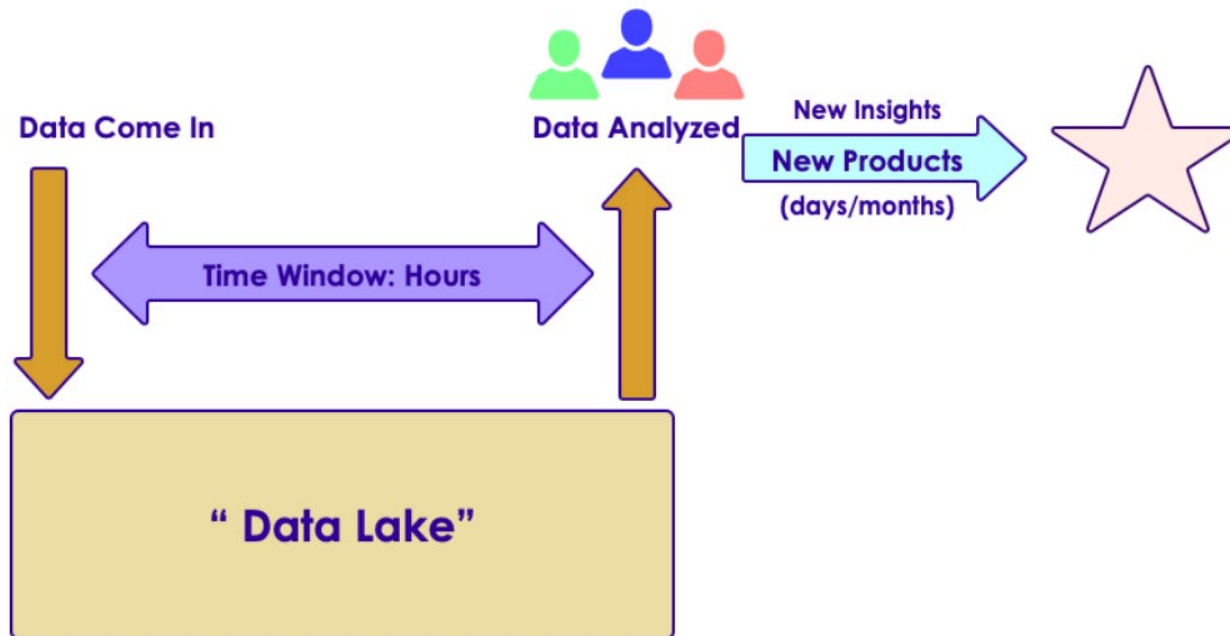


*Presents*

# **Streaming with Kafka**


# Big Data Evolution: Version 1

- ▶ Decision times: batch ( hours / days)
  - ✓ Use cases: Modeling, ETL, Reporting



# Moving Towards Fast Data: Version2

- ▶ Decision time: (near) real time
  - ✓ Seconds (or milliseconds)
- ▶ Use Cases
  - ✓ Alerts (medical/security)
  - ✓ Fraud detection
- ▶ Streaming is becoming more prevalent
  - ✓ Connected Devices
  - ✓ Internet of Things
- ▶ Beyond Batch
  - ✓ We need faster processing and analytics



# Streaming Use Cases

- ▶ Netflix
  - ✓ Recommendations 450 billion events/day
- ▶ Weather Company
  - ✓ Analyze weather sensor data
  - ✓ Billions of events/day
  - ✓ Multi-Petabyte (PB) traffic daily



# Real Time / Near Real Time

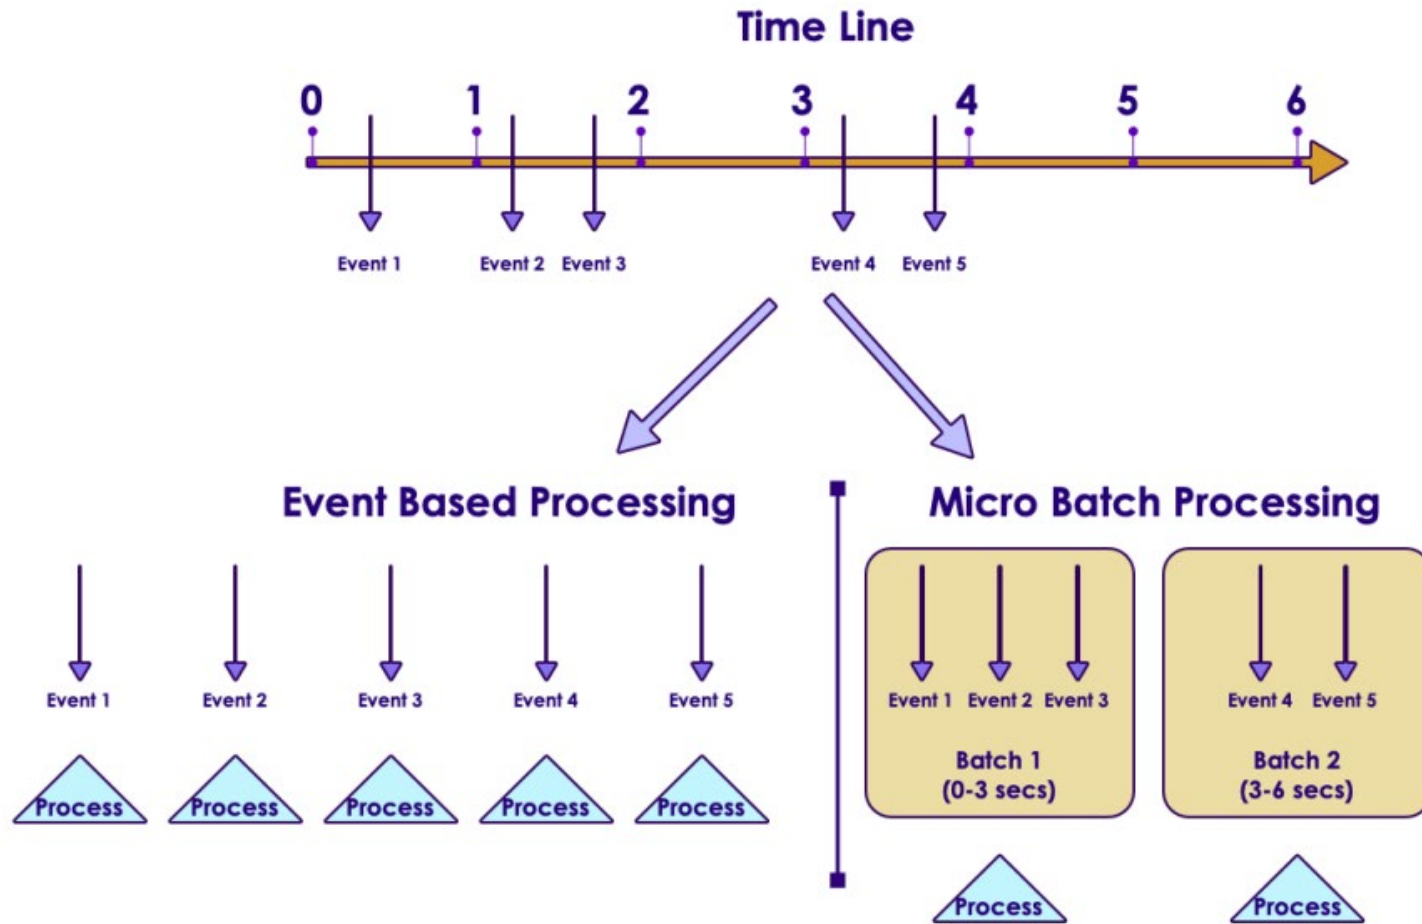
- ▶ The 'real' real time is in milliseconds order
  - ✓ DB query returns in 2 ms
- ▶ 'Near real time' is seconds
  - ✓ We can process an event within 3 seconds of its generation time

Name	Time	Example
Hard real time	Single order ms, sub milli seconds 1 ms, 0.5 ms	Space shuttle control systems
Credit card transaction processing	50 ms, 300 ms	Db queries
Sending Emails	2 secs +	Stream processing latency
	1 min +	Mini batch queries

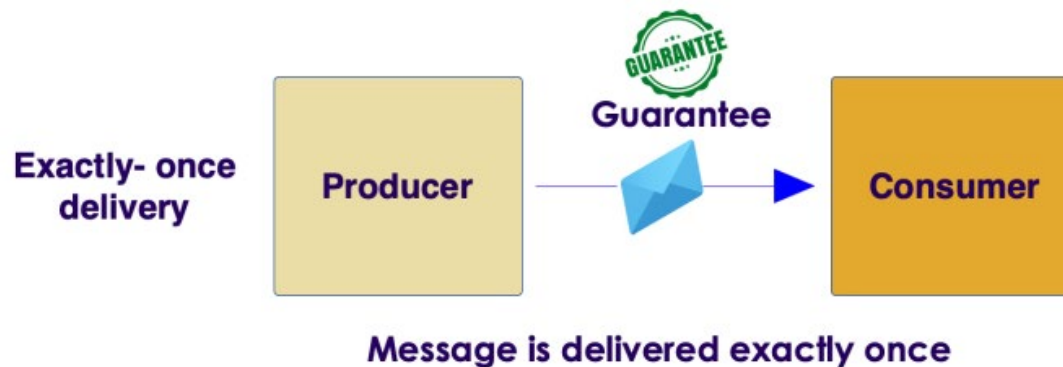
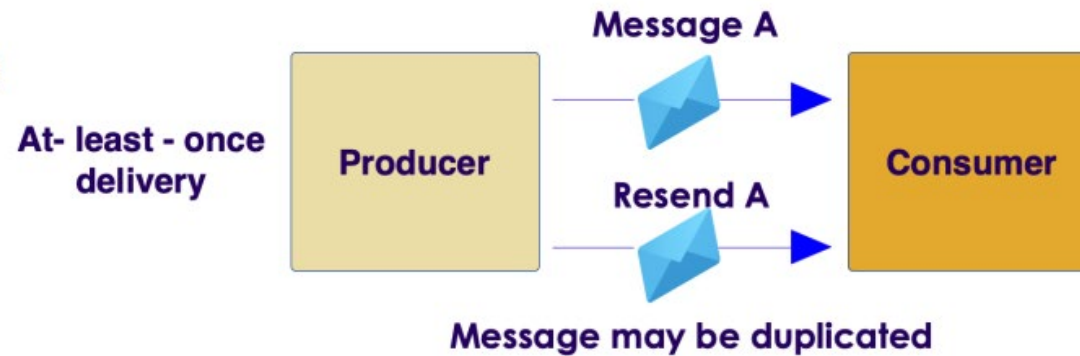
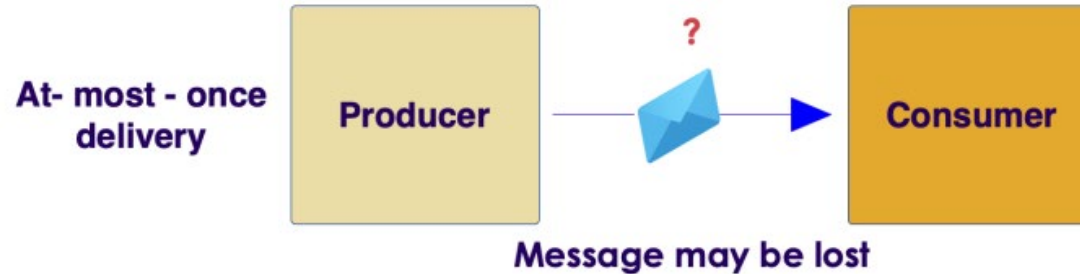
# Streaming Concepts

- ▶ Processing model
  - ✓ Event based or micro batch based
- ▶ Processing guarantees
  - ✓ At least once
  - ✓ At most once
  - ✓ Exactly once
- ▶ State management
  - ✓ Event time vs. Arrival time
- ▶ Window Operations
- ▶ Back-pressure adjustment

# Event Based Vs. Batch



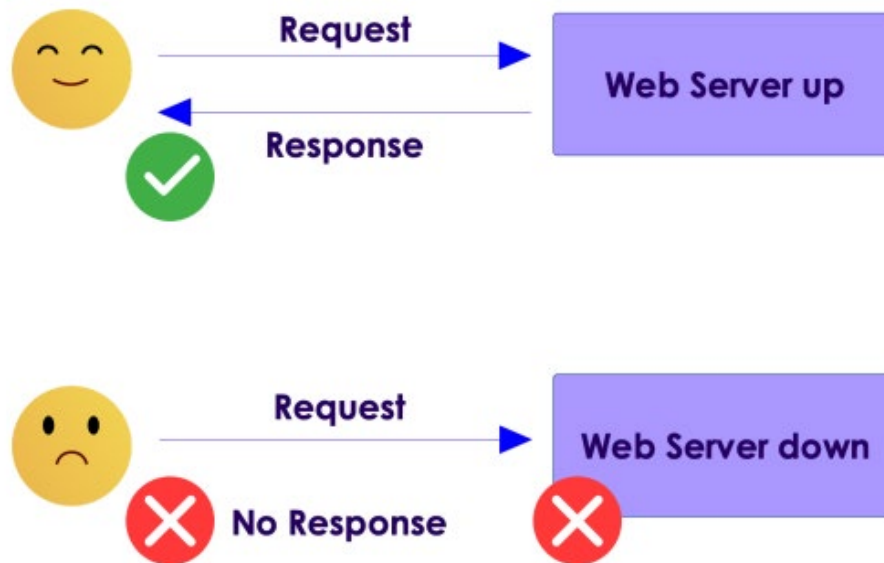
# Processing Guarantees





# At Most Once

- ▶ Event is sent only once
  - ✓ No duplicate processing
  - ✓ Events can be dropped due to crashes or heavy load
  - ✓ E.g. Web requests (if the web server is busy, requests are dropped)



# At Least Once


- ▶ All events are guaranteed to be processed (no dropped events)
  - ✓ However, events can be processed more than once
  - ✓ In case of failure recovery, events can be re-sent and processed again.
- ▶ Most common implementation
  - ✓ Frameworks: All (Storm, Spark, NiFi, Samza, Flink)

# Handling Duplicate Events

- ▶ A resilient streaming system, has to be ready to handle duplicate events
- ▶ We have 2 scenarios:
  - ✓ First one, we are inserting a new record for each event received. This will result in duplicate records in the database
  - ✓ Second one, we are checking to see if the event is processed already, only if not, then a new record is inserted
- ▶ Second approach is more resilient, can deal with duplicate events
  - ✓ This is called idempotent processing (no side effects for duplicate events)

# Handling Duplicate Events

Scenario-1: Duplicate records created



Scenario-2: No Duplicate records created



# Exactly Once

- ▶ Events are guaranteed to be processed once and only once
  - ✓ No dropped events
  - ✓ No duplicate processing
  - ✓ Frameworks: Storm (with Trident), Flink, Spark, Samza
- ▶ Sample applications
  - ✓ Credit card processing



**Mathias Verraes**  
@mathiasverraes

Follow

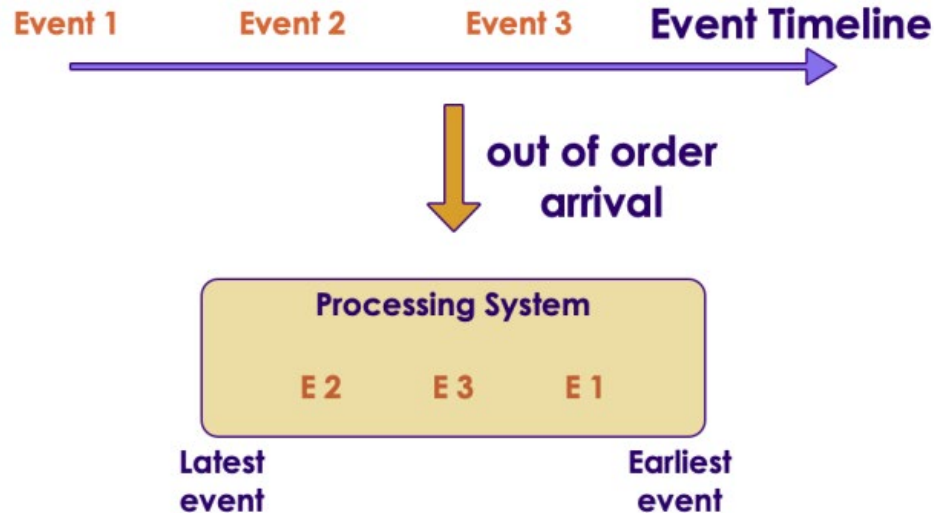


There are only two hard problems in distributed systems:  
1. Guaranteed order of messages  
2. Exactly-once delivery

11:40 AM - 14 Aug 2015


# Event Time and Arrival Time

- ▶ Event Time: When the event occurred / generated
- ▶ Arrival Time: When event arrives for processing
- ▶ Event Time < Arrival Time
  - ✓ Some times events may arrive 'out of order' (due to network lag, outage ..etc)



# 3 Tier Streaming Architecture

- ▶ Here is a simplified streaming architecture
- ▶ We see 3 distinct stages
  - ✓ Ingest stage captures data
  - ✓ Processing handles the data
  - ✓ And the processed data is stored in Storage layer



# Ingest / Capture

- ▶ This layer:
  - ✓ Captures incoming data
  - ✓ Acts as a 'buffer' - smoothes out bursts So even if our processing offline, we won't loose data
- ▶ Choices
  - ✓ **Kafka**
  - ✓ Queues (MQ, JMS ..etc)
  - ✓ Cloud based queues like Amazon Kinesis



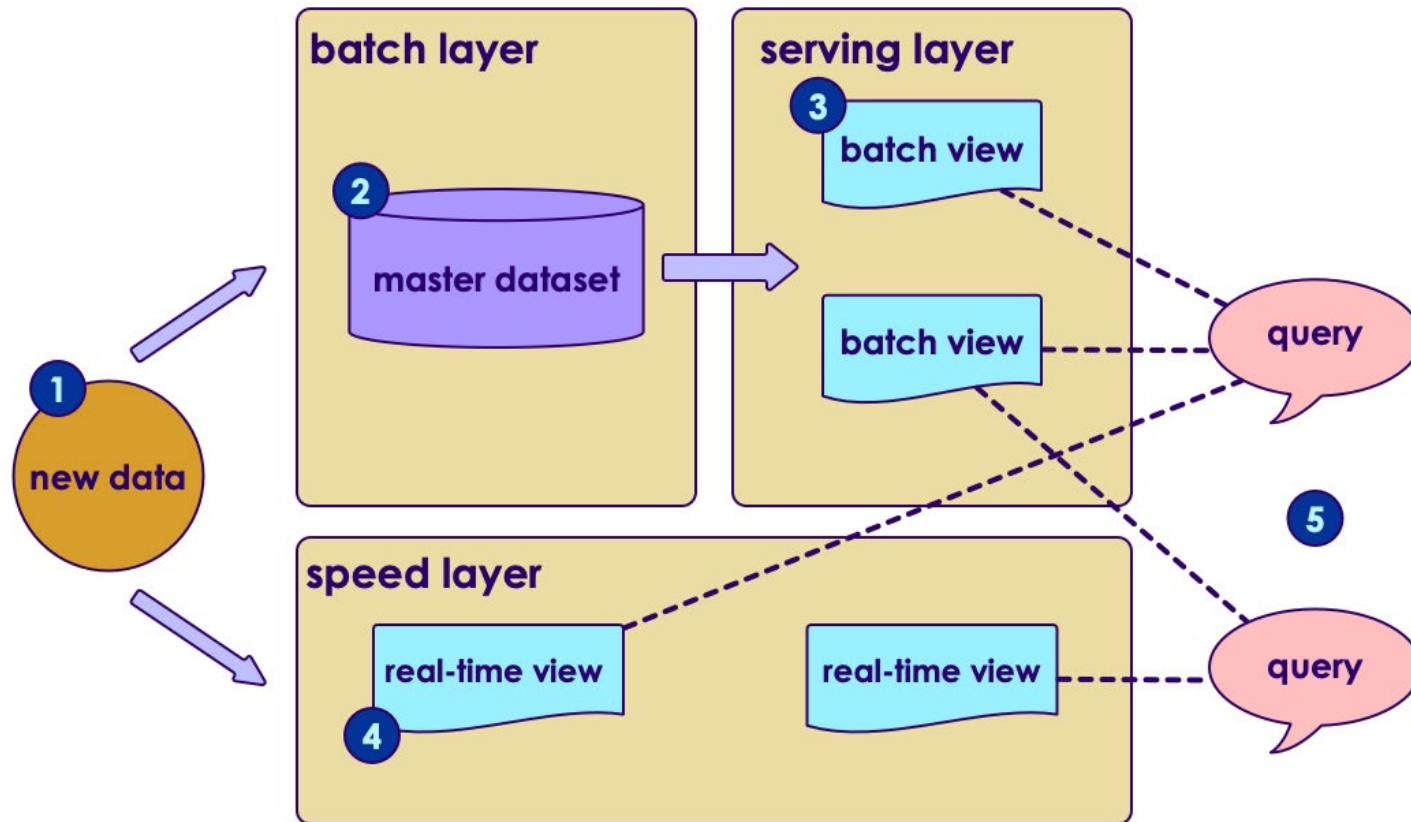
# Processing

- ▶ We need to process events with low latency
  - ✓ (milliseconds to seconds)
- ▶ There are many stream/event processing frameworks available
  - ✓ Storm
  - ✓ Spark
  - ✓ NiFi
  - ✓ Samza
  - ✓ Flink
  - ✓ Beam

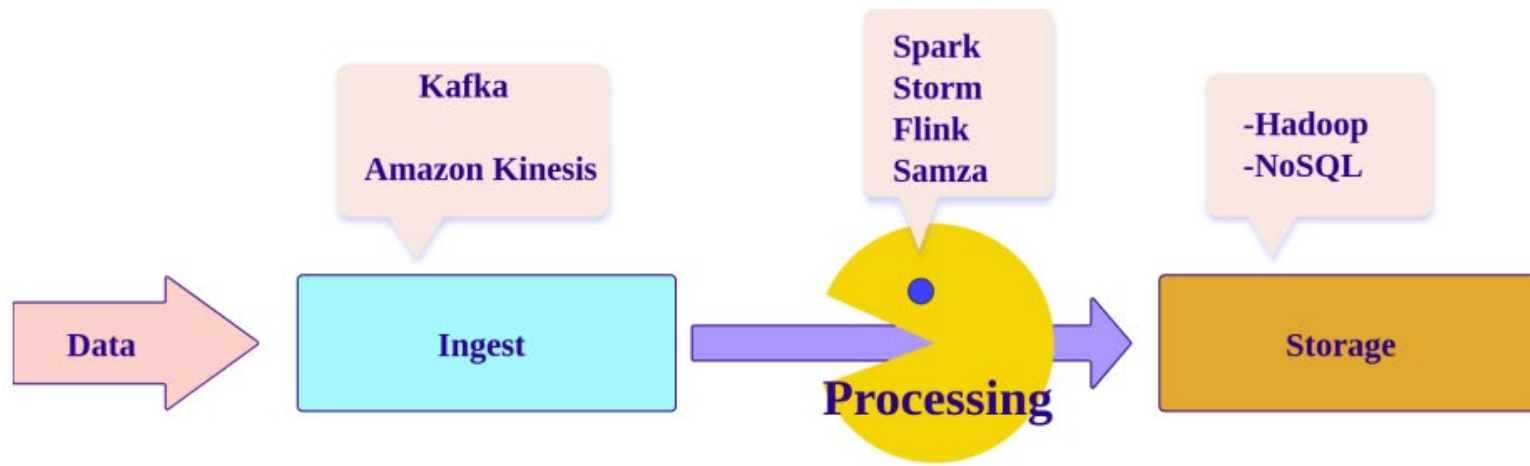
# Storage

- ▶ After processing, they are stored for later retrieval
- ▶ Two choices:
  - ✓ Real time store
  - ✓ 'Forever' store
- ▶ Real Time Store
  - ✓ Need to absorb data in real time
  - ✓ Usually a NoSQL storage (HBase, Cassandra ...etc)
  - ✓ May contain subset of data (last 1 year ..etc)
- ▶ 'Forever store'
  - ✓ Needs to store massive amounts of data
  - ✓ Support analytics (usually batch)
  - ✓ Hadoop / HDFS

# Lambda Architecture



# Streaming Stack - Summary




# Apache Kafka

- ▶ Kafka is a Publisher / Subscriber (Pub-Sub) messaging system
  - ✓ Distributed
    - Scales seamlessly
  - ✓ High throughput
    - Capable of handling billions of messages per day
  - ✓ Replicated
    - Safeguards data in case of machine failures
- ▶ Created @ LinkedIn in 2010
  - ✓ Now Apache Project (Open Source)



# Why Is Kafka Very Fast?

- ▶ Write: Disk writes are buffered in page cache
- ▶ Read: The data from page cache can be transferred to network interface very efficiently
- ▶ 99% of the time data is read from page cache, no disk access at all



# Kafka Features

Feature	Kafka	Other Queue Systems
Deleting messages	Clients can not delete. ,Kafka auto-expires messages	Clients can delete
Message processing order	Can read in or out-of order	Usually read in order
Message processing guarantee	Kafka guarantee no duplicate processing of a message	Usually no
Concurrent read / write	Supported.,High throughput	Low throughput due to locking & blocking
Message priorities	None	Yes
Message ACKs,(Client notify producer that a message is processed)	No	May be

# Questions

