



*Presents*

# **Kafka Architecture**

# Kafka Command Utilities in BIN

- ▶ Starting Kafka brokers
  - ✓ bin/kafka-server-start
  - ✓ bin/kafka-server-stop
- ▶ Managing topics
  - ✓ bin/kafka-topics: Lists / create / delete topics
- ▶ Sending Messages
  - ✓ bin/kafka-console-producer.sh
- ▶ Consuming messages
  - ✓ bin/kafka-console-consumer.sh

# Creating Topics

```
$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --list
# ... empty ...

## create a topic with one replica and two partitions
$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --create
--topic test --replication-factor 1 --partitions 2

$ bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic test

# Topic:test      PartitionCount:2      ReplicationFactor:1      Configs:
# Topic: test      Partition: 0      Leader: 0      Replicas: 0      Isr: 0
# Topic: test      Partition: 1      Leader: 0      Replicas: 0      Isr: 0
```

# Using Producer / Consumer Utils

- ▶ bin/ kafka-console-producer:
  - ✓ utility for producing messages
- ▶ bin/kafka-console-consumer:
  - ✓ utility for reading messages

```
$ bin/kafka-console-producer.sh  
--broker-list localhost:9092  
--topic test  
  
Hello  
World  
Goodbye  
world
```

Producer:  
Type input here



```
$ bin/kafka-console-consumer.sh  
--bootstrap-server localhost:9092  
--topic test  
  
Hello  
World  
Goodbye  
world
```

Consumer:  
Data will be read from Kafka  
and will show up here

# Kafka Clients

- ▶ Java is the 'first class' citizen in Kafka
  - ✓ Officially maintained
- ▶ Python on par with Java
  - ✓ Maintained by Confluent.io
- ▶ Other language libraries are independently developed
  - ✓ May not have 100% coverage
  - ✓ May not be compatible with latest versions of Kafka

# Kafka Java API

- ▶ Rich library that provides high level abstractions
  - ✓ No need to worry about networking / data format ..etc
- ▶ Write message / Read message
- ▶ Supports native data types
  - ✓ String
  - ✓ Bytes
  - ✓ Primitives (int, long ...etc.)

# Java Producer Code (Abbreviated)

```
// ** 1 **
import java.util.Properties;
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;
import org.apache.kafka.clients.producer.ProducerConfig;
import org.apache.kafka.common.serialization.StringSerializer;
import org.apache.kafka.common.serialization.IntegerSerializer;
...

// ** 2 **
Properties props = new Properties();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.CLIENT_ID_CONFIG, "SimpleProducer");
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, IntegerSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());

KafkaProducer< Integer, String > producer = new KafkaProducer<>(props);

// ** 3 **
String topic = "test";
Integer key = new Integer(1);
String value = "Hello world";
ProducerRecord< Integer, String > record = new ProducerRecord<>(topic, key, value);
producer.send(record);
producer.close();
```

# Producer Code Walkthrough

```
// ** 2 ** Recommended approach: use constants
```

```
import org.apache.kafka.clients.producer.ProducerConfig;
import org.apache.kafka.common.serialization.StringSerializer;
import org.apache.kafka.common.serialization.IntegerSerializer

Properties props = new Properties();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.CLIENT_ID_CONFIG, "SimpleProducer");
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, IntegerSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());

KafkaProducer < Integer, String > producer = new KafkaProducer<>(props);
```

```
// ** 2 ** another approach
```

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("client.id", "SimpleProducer");
props.put("key.serializer", "org.apache.kafka.common.serialization.IntegerSerializer");
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

KafkaProducer < Integer, String > producer = new KafkaProducer<>(props);
```



# Producer Code Walkthrough

- ▶ Each record represents a message
  - ✓ Here we have a <key,value> message
  - ✓ send() doesn't wait for confirmation
- ▶ We send in batches
  - ✓ For increased throughput
  - ✓ Minimize network round trips

```
// ** 3 **  
String topic = "test";  
Integer key = new Integer(1);  
String value = "Hello world";  
ProducerRecord< Integer, String > record = new ProducerRecord<> (topic, key, value);  
producer.send(record);  
producer.close();
```

# Producer Properties

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("client.id", "SimpleProducer");
props.put("acks", "all");
props.put("retries", 0);
props.put("batch.size", 16384); // 16k
props.put("linger.ms", 1);
props.put("buffer.memory", 33554432); // 32 M
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, IntegerSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());

KafkaProducer < Integer, String > producer = new KafkaProducer<>(props);


for(int i = 0; i < 100; i++) {
    producer.send(new ProducerRecord < String, String >(
        "my-topic", Integer.toString(i), Integer.toString(i)));
}
producer.close();
```

# Producer Acknowledgements



ACK	Description	Speed	Data safety
acks=0	<ul style="list-style-type: none"><li>- Producer doesn't wait for any acks from broker,</li><li>- Producer won't know of any errors</li></ul>	High	<p>Low</p> <p>No guarantee that broker received the message</p>
acks=1, (default)	<ul style="list-style-type: none"><li>- Broker will write the message to local log,</li><li>- Does not wait for replicas to complete</li></ul>	Medium	<p>Medium</p> <p>Message is at least persisted on lead broker</p>
acks=all	<ul style="list-style-type: none"><li>- Message is persisted on lead broker and in replicas,</li><li>- Lead broker will wait for in-sync replicas to acknowledge the write</li></ul>	Low	<p>High</p> <p>Message is persisted in multiple brokers</p>

# Producer Acknowledgements



# Consumer Code (Abbreviated)

```
import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.clients.consumer.ConsumerConfig;
import org.apache.kafka.common.serialization.StringDeserializer;
import org.apache.kafka.common.serialization.IntegerDeSerializer

...

Properties props = new Properties(); // ** 1 **
props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ConsumerConfig.GROUP_ID_CONFIG, "group1");
props.put(ConsumerConfig.CLIENT_ID_CONFIG, "Simple Consumer");
props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, IntegerDeSerializer.class.getName());
props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, StringDeserializer.class.getName());

KafkaConsumer < Integer, String > consumer = new KafkaConsumer<>(props);

consumer.subscribe(Arrays.asList("topic1")); // ** 2 **

try {
    while (true) {
        ConsumerRecords < Integer, String > records = consumer.poll(Duration.ofMillis(1000)); // ** 3 **
        System.out.println("Got " + records.count() + " messages");
        for (ConsumerRecord < Integer, String > record : records) {
            System.out.println("Received message : " + record);
        }
    }
}
finally {
    consumer.close(Duration.OfSeconds(60));
}
```

# Consumer Code Walkthrough

```
Properties props = new Properties(); // ** 1 **
props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ConsumerConfig.GROUP_ID_CONFIG, "group1");
props.put(ConsumerConfig.CLIENT_ID_CONFIG, "Simple Consumer");
props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, IntegerDeSerializer.class.getName());
props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, StringDeSerializer.class.getName());

KafkaConsumer < Integer, String > consumer = new KafkaConsumer<>(props);

consumer.subscribe(Arrays.asList("topic1")); // ** 2 **
```

- ▶ bootstrap.servers: "broker1:9092,broker2:9092"
  - ✓ Connect to multiple brokers to avoid single point of failure
  - ✓ group.id: consumers belong in a Consumer Group
  - ✓ We are using standard serializers
- ▶ Consumers can subscribe to one or more subjects // \*\* 2 \*\*


# Consumer Code Walkthrough

```
try {  
    while (true) {  
        ConsumerRecords < Integer, String > records = consumer.poll(Duration.ofMillis(1000)); // ** 3 **  
        System.out.println("Got " + records.count() + " messages");  
        for (ConsumerRecord < Integer, String > record : records) {  
            System.out.println("Received message : " + record);  
        }  
    }  
}  
finally {  
    consumer.close();  
}
```

- ▶ Consumers must subscribe to topics before starting polling
  - ✓ `Consumer.subscribe ("test.*")` // wildcard subscribe
  - ✓ Poll: This call will return in 1000 ms, with or without records
  - ✓ Must keep polling, otherwise consumer is deemed dead and the partition is handed off to another consumer

# Consumer Poll Loop

- ▶ Polling is usually done in an infinite loop.
  - ✓ First time poll is called
  - ✓ Finds the GroupCoordinator
  - ✓ Joining Consumer Group
  - ✓ Receiving partition assignment
- ▶ Work done in poll loop
  - ✓ Usually involves some processing
  - ✓ Saving data to a store
  - ✓ Don't do high latency work between polls; otherwise the consumer could be deemed dead.
  - ✓ Do heavy lifting in a separate thread





# ConsumerRecord

- ▶ org.apache.kafka.clients.consumer.ConsumerRecord  
<K,V>
- ▶ K key(): key for record (type K), can be null
- ▶ V value(): record value (type V - String / Integer ..etc)
- ▶ String topic(): Topic where this record came from
- ▶ int partition(): partition number
- ▶ long offset(): long offset in

```
ConsumerRecords < Integer, String > records = consumer.poll(Duration.ofMillis(1000));  
for (ConsumerRecord < Integer, String > record : records) {  
    System.out.printf("topic = %s, partition = %d, offset = %d,  
        key= %s, value = %s\n",  
        record.topic(), record.partition(), record.offset(),  
        record.key(), record.value());  
}
```

# Configuring Consumers

- ▶ `max.partition.fetch.bytes` (default : 1048576 (1M))
  - ✓ Max message size to fetch. Also see `message.max.bytes` broker config
- ▶ `session.timeout.ms` (default : 30000 (30 secs))
  - ✓ If no heartbeats are not received by this window, consumer will be deemed dead and a partition rebalance will be triggered

```
Properties props = new Properties(); // ** 1 **  
  
...  
props.put("session.timeout.ms", 30000); // 30 secs  
props.put("max.partition.fetch.bytes", 5 * 1024 * 1024); // 5 M  
  
KafkaConsumer < Integer, String > consumer = new KafkaConsumer<>(props);
```

# Clean Shutdown Of Consumers

- ▶ Consumers poll in a tight, infinite loop
- ▶ Call ' consumer.wakeup () ' from another thread
- ▶ This will cause the poll loop to exit with ' WakeupException '

```
try {
    while (true) {
        ConsumerRecords < Integer, String > records = consumer.poll(100);
        // handle events
    }
}
catch (WakeupException ex) {
    // no special handling needed, just exit the poll loop
}
finally {
    // close will commit the offsets
    consumer.close();
}
```

# Signaling Consumer To Shutdown

- ▶ Can be done from another thread or shutdown hook
- ▶ 'consumer.wakeup ()' is safe to call from another thread

```
Runtime.getRuntime().addShutdownHook(new Thread() {  
    public void run() {  
        System.out.println("Starting exit...");  
        consumer.wakeup(); // signal poll loop to exit  
        try {  
            mainThread.join(); // wait for threads to shutdown  
        } catch (InterruptedException e) {  
            e.printStackTrace();  
        }  
    }  
})
```

# Questions

